Stance Detection for the Fake News Challenge Dataset using Deep Learning

Anak Agung Ngurah Bagus Trihatmaja and Shishir Kurhade

Abstract

The problem of fake news has arisen recently as a threat to high-quality journalism and well-informed public disclosure. The goal of fake news challenge is to explore how artificial intelligence technologies, particularly machine learning and natural language processing, might be leveraged to combat the fake news problem[1]. The goal of our project is to develop machine learning models to predict a stance label (Agrees, Disagrees, Related, Unrelated) with respect to the title for a respective news article. For this purpose, we will use the gated recurrent unit (GRU) to predict the labels. As a baseline to measure performance we will also solve the problem using logistic regression method.

1 Introduction

In this project, we try to combat a serious problem in our media using machine learning techniques. In a poll conducted by Pew Research Center, 64% of US adults said that fake news has caused a great deal of confusion about the basic facts of current issues and events[2]. This problem views the task of fake-news detection as a stance detection problem which is a labeling task. We want to automatically classify a news into four labels, which are unrelated, agrees, disagrees, and discusses.

A reasoning for these labels is as follows:

1. **Agrees**: The body text agrees with the headline.

2. **Disagrees**: The body text disagrees with the headline.

3. **Discusses**: The body text discuss the same topic as the headline, but does not take a position

4. **Unrelated**: The body text discusses a different topic than the headline

The classifier that we build could later be used as a base of a fake news detection tool that can automatically categorize the news into the stances given.

2 Technical Approach

Our problem of stance detection can be viewed as that of text classification with each of the stances as a possible class. We use two approaches to solve this problem viz. logistic regression and deep neural networks.

2.1 Logistic Regression

Logistic is the appropriate regression analysis to conduct when the dependent variable is dichotomous (binary). Logistic regression is used to describe data and to explain the relationship between one dependent binary variable (stances in our case) and one or more independent variables (words in the vocabulary). In our problem setting we aim at classifying every news article body in a stance with respect to the title. For representing the text as features we use a representation form called tf-idf (term frequency inverse document frequency) matrix. Here, each document is represented as a vector with respect to the vocabulary of words. The model concatenates the text from the title and the body of an news article to prepare a vectorised representation. This representation increases the values of weights for words proportionally to the number of times a word appears in the document and is offset by the frequency of the word in the corpus. This is helpful in our task as the news articles are of variable length, thus giving us a robust feature representation. We perform one vs all regression for each stance. In pseudocode,

the training algorithm for an One vs All learner constructed from a binary classification learner L is as follows:

Algorithm 1 Logistic regression

Input
    L (training algorithm for binary classifiers)
    samples X
    labels y where $y_i \in \{1, ..., K\}$ is the label for the sample X_i
Output
    a list of classifiers f_k for $k \in \{1, ..., K\}$

  1: **procedure** LOGISTICREGRESSION
  2:    **for** $k \leftarrow 1, K$ **do**
  3:       **if** $y_i = k$ **then**
  4:          $z_i \leftarrow 1$
  5:       **else**
  6:          $z_i \leftarrow 0$
  7:       **end if**
  8:       $f_k \leftarrow L(X, z)$ ▷ Apply L to X, z to obtain f_k
  9:    **end for**
10: **end procedure**

    Making decisions means applying all classifiers to an unseen sample x and predicting the label k for which the corresponding classifier reports the highest confidence score f_k.

2.2   Deep Neural Networks

Neural networks are a set of algorithms, modelled loosely after the human brain, that are designed to recognize patterns. They interpret sensory data through a kind of machine perception, labeling or clustering raw input. The patterns they recognize are numerical, contained in vectors, into which all real-world data, be it images, sound, text or time series, must be translated. Traditionally RNNs have been widely used in NLP tasks because of their memory capabilities helping them to retain the sequence information in the text, thus developing better language models. However, they suffer from problems of vanishing gradient and variants of RNNs consisting of LSTM and GRU cells solve this problem. As GRUs are faster and computationally less expensive we implemented our network using GRU cells in Tensorflow to predict the stances. For feature representation we embed each word in the title and article body using GloVe representations into vectors of size 50. These are pre-trained word vector representations made readily available by Stanford NLP group. We have tried two variants of neural networks for our task:

1. Vanilla RNN with single hidden layer

2. RNN with multiple hidden layers and dropout

    The input layer consisting of 50 GRU cells accepts one word at a time from the input news title and article whose output is then used as an input for the next hidden layer. This pipelining of outputs as input to next layer continues for all hidden layers and the final layer which consists of softmax layer classifies the stance. We use the cross entropy loss function to calculate the loss and optimize it using the Adam Optimizer. It is important to note that the neural network preserves the sequential information present in the sentences to generate weights for classifying the test samples unlike the baseline method which is a bag of words model.

3   Experimental Results

3.1   Dataset Overview

The data provided by the Fake News Challenge consists of headline, body, and stace. For training, there are two csv files:

1. *train_bodies.csv*: contains the body text of articles with its ID

2. *train_stances.csv*: contains labeled stances for pairs of article headlines and article bodies, in which the article bodies refer to the bodies in train_bodies.csv

The distribution of the data is as follows:

Rows	Unrelated	Discuss	Agree	Disagree
49972	0.73131	0.17828	0.0736012	0.0168094

Table 1: Dataset distribution

We will roughly use 4000 samples as our development set, for choosing hyperparameter and performance evaluation, and use the rest for training. Meanwhile, for the test dataset we use 4000 samples for testing our model.

3.2 Data Pre-processing

For both logistic and RNN, we will do the same data cleaning procedures. There is a difference in feature representation for each of the models. The data cleaning procedures we perform is as follows:

1. Normalize the case

2. Remove all punctuations and non-alphabetic symbols

3. Replace the shorten words with the normal words, for example: you'll into you will

In case of logistic regression to reduce dimentionality, we perform lemmatization on the text generated from concatinating headlines and the bodies. We then build our tf-idf matrix with this vocabulary and data using scikit-learn [3]. The tf-idf matrix looks like below.

	Word 1	Word 2	Word 3	Word 4	...
Document 1	0.01	0.0	0.05	0.13	0.12
Document 2	0.2	0.0	0.04	0.4	0.1
Document 3	0.12	0.0	0.101	0.012	0.0

Table 2: Input representation for logistic regression

Each row in the table represents one document we have in the dataset. While, each column represents a word present in the vocabulary.

For RNN using GRUs, we represents our data as word vectors. The text from the corpus will be mapped to corresponding vectorized forms using pre-trained *GloVe* representations–freely available on the Stanford NLP group website[4, 5]. Since the text sequences observed will be of variable length we will pad all sequences to the length of the maximum length text sequence before inputting it to our model. The representation of the matrix can be seen as follows-

	Headlines	Bodies
Document 1	[0.23, 0.12, ...], [0.123, 0.14, ...], ...	[0.05, 0.012, ...], [0.121, 0.123, ...]
Document 2	[0.23, 0.12, ...], [0.119, 0.4, ...], ...	[0.11, 0.011, ...], [0.109, 0.89, ...]
Document 3	[0.28, 0.1, ...], [0.134, 0.3, ...], ...	[0.05, 0.001, ...], [0.121, 0.189, ...]

Table 3: Input representation for neural network

As our training dataset is skewed and has a majority of stances as 'unrelated', we separate our training dataset into two. The first dataset contains related vs related stances where (agree, disagree, discuss) are together bucketed into 'related'. The second is a dataset containing only agree, disagree, discuss. We then train our models separately for each dataset.

In the case of RNNs, we represent the label as a vector. For the unrelated and related stance, we represent it as $[1, 0]$ if it is unrelated and $[0, 1]$ as related. While, for the second dataset we label our dataset as $[1, 0, 0]$ if the stance is agree, $[0, 1, 0]$ if the stance is disagree and $[0, 0, 1]$ as discuss. The final structure of our data for RNN looks as follows.

3

	Headlines	Bodies	Stances
Document 1	[0.23, 0.12, ...], [0.123, 0.14, ...], ...	[0.05, 0.012, ...], [0.121, 0.123, ...]	[0,1]
Document 2	[0.23, 0.12, ...], [0.119, 0.4, ...], ...	[0.11, 0.011, ...], [0.109, 0.89, ...]	[1, 0]
Document 3	[0.28, 0.1, ...], [0.134, 0.3, ...], ...	[0.05, 0.001, ...], [0.121, 0.189, ...]	[1,0]

Table 4: Final input representation for neural network

3.3 Experimental Settings

Our dataset consists of 4 classes (unrelated, agrees, disagrees, discuss) with headline and body as features for the RNNs. Each word in GRU is represented as a vector of length 50. For the logistic regression, we have 33795 unique and lemmatized words used as features represented in form of tf-idf matrix.

We train our models on a computer with Intel i7-8550U CPU @ 1.80GHz 8 and 16 GB of RAM. The neural network takes four hours to train while the logistic regression model takes about ten minutes For getting the test result, we run our test for the neural network on AWS (Amazon Web Services) EC2 p3.2xlarge package running Ubuntu with specifications as-

GPUs - Tesla V100	GPU Memory (GB)	vCPUs	Memory (GB)
1	16	8	61

Table 5: AWS specification

3.4 Validation

We get the best value for regularization using cross validation on our logistic regression model. We run the the logistic with several C values. The C in denotes the inverse of regularization strength. The results of the cross validation are as follows-


Figure 1: Test validation scores against value of C for each dataset

	C = 1e-15	C = 1	C = 1000
Agree	0 %	71 %	68 %
Disagree	0 %	0 %	46 %
Discuss	59 %	77 %	79 %

(a) Precision score for different value of C in related stances dataset

	C = 1e-15	C = 1	C = 1000
Agree	0 %	63 %	66 %
Disagree	0 %	0 %	19 %
Discuss	100 %	89 %	85 %

(b) Recall score for different value of C in related stances dataset

For the GRU, we run the validation using different number of neurons. We then pick the best result.

Figure 2: Test validation scores against number of hidden neurons for each dataset

The graph above shows the effect of number of neurons in a layer for each dataset. We can see that the accuracy reaches its peak when there are 60 neurons in hidden layer for unrelated vs related dataset, while for the related stances the accuracy is quite similar for every number of neurons. Therefore we choose 60 as the number of neurons in the hidden layers.

3.5   Results and Analysis

Accuracy	Logistic regression	RNN with 1 hidden layer	RNN with 2 hidden layers
Related vs unrelated	52.76 %	81.58 %	78.66 %
Related stances	66.03 %	56.94 %	61.19 %

Table 7: Results from RNN and logistic regression

We also calculated the accuracies for three hidden layers in RNN and the accuracy for related stances dropped drastically to 21.2 %. For related vs unrelated stances (agree, disagree, discuss together form unrelated) we use the entire dataset while training and for the related stances training we only use agree, disagree, discuss stances which form around 27 % of the total training data. Logistic regression performs better on smaller datasets and hence has high accuracy for related stances. The neural networks give a good performance when there is a large dataset available to train them. This explains their higher accuracy as compared to logistic regression for related vs unrelated stances. Considering the overall performances of the models we conclude that a RNN model with two hidden layers is best suitable for our problem.

4   Participants Contribution

Tasks	Trihatmaja	Kurhade
Data pre-processing for logistic regression		✓
Data pre-processing for RNN	✓	
RNN implementation	✓	
Logistic regression implementation		✓
Server environment setup	✓	
RNN improvement		✓
Logistic regression cross validation	✓	
RNN test		✓
Presentation, project proposal and report	✓	✓

References

[1] Dean Pomerleau and Delip Rao. Fake news challenge.

[2] Michael Barthel, Amy Mitchell, and Jesse Holcomb. Many americans believe fake news is sowing confusion, Dec 2016.

[3] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.

[4] Steven Bird, Ewan Klein, and Edward Loper. *Natural Language Processing with Python*. O'Reilly Media, Inc., 1st edition, 2009.

[5] Jeffrey Pennington, Richard Socher, and Christopher D Manning. Glove: Global vectors for word representation. In *EMNLP*, volume 14, pages 1532–1543, 2014.